

# Linear.Model

Yue Wen

August 11, 2017

## 1. Missing data imputation

I will summarize what I did:

- 1) For tax/area related missing data, because there are relationship between different tax featres, it works the same way for the area feature. Therefore, we used the related feature to impute the missing data by simple calculation.
- 2) For zip code/city, what we did here is to use longitude/latitude to find its nearest neighbor and them use the same zip/city.
- 3) For the missing features that contains more then 80% NA, delete them.
- 4) For categorical features that contains a lot of NA, create new level missing.
- 5) For total area, using library(mice) to impute.

Because it takes long to run the imputation part, I save the result to the csv, and we don't evaluate the code below anymore. Instead, we directly import the data generated.

```
# 1. read/prepare data
#####
setwd("D:/data_camp/zillow_project")
train <- read.csv('train_property.csv', stringsAsFactors = F)
train$trans_year <- sapply(strsplit(train$transactiondate, '-'), '[[', 1)
train$trans_month <- sapply(strsplit(train$transactiondate, '-'), '[[', 2)
train$trans_day <- sapply(strsplit(train$transactiondate, '-'), '[[', 3)
train$trans_weekday <- weekdays(as.Date(train$transactiondate))
train$trans_DATE <- as.Date(train$transactiondate)
# train rename
train <- plyr::rename(train,
                      c("parcelid"="id_parcel",
                        "transactiondate" = "trans_date",
                        "yearbuilt" = "build_year",
                        "basementsqft"="area_base_living",
                        "yardbuildingsqft17"="area_patio",
                        "yardbuildingsqft26"="area_shed",
                        "poolsum"="area_pool",
                        "lotsizesquarefeet"="area_lot",
                        "garagetotalsqft"="area_garage",
                        "finishedfloor1squarefeet" = "area_firstfloor_finished",
                        "calculatedfinishedsquarefeet" = "area_total_calc",
                        "finishedsquarefeet6" = "area_base",
                        "finishedsquarefeet12" = "area_live_finished",
                        "finishedsquarefeet13" = "area_liveperi_finished",
                        "finishedsquarefeet15" = "area_total_finished",
                        "finishedsquarefeet50" = "area_unknown",
                        "unitcnt" = "num_unit",
                        "numberofstories" = "num_story",
                        "roomcnt" = "num_room",
```

```

    "bathroomcnt" = "num_bathroom",
    "bedroomcnt" = "num_bedroom",
    "calculatedbathnbr" = "num_bathroom_calc",
    "fullbathcnt" = "num_bath",
    "threequarterbathnbr" = "num_75_bath",
    "fireplacecnt" = "num_fireplace",
    "poolcnt" = "num_pool",
    "garagecarcnt" = "num_garage",
    "regionidcounty" = "region_county",
    "regionidcity" = "region_city",
    "regionidzip" = "region_zip",
    "regionidneighborhood" = "region_neighbor",
    "taxvaluedollarcnt" = "tax_total",
    "structuretaxvaluedollarcnt" = "tax_building",
    "landtaxvaluedollarcnt" = "tax_land",
    "taxamount" = "tax_property",
    "assessmentyear" = "tax_year",
    "taxdelinquencyflag" = "tax_delinquency",
    "taxdelinquencyyear" = "tax_delinquency_year",
    "propertyzoningdesc" = "zoning_property",
    "propertylandusetypeid" = "zoning_landuse",
    "propertycountylandusecode" = "zoning_landuse_county",
    "fireplaceflag" = "flag_fireplace",
    "hashottuborspa" = "flag_tub",
    "buildingqualitytypeid" = "quality",
    "buildingclasstypeid" = "framing",
    "typeconstructiontypeid" = "material",
    "decktypeid" = "deck",
    "storytypeid" = "story",
    "heatingorsystemtypeid" = "heating",
    "airconditioningtypeid" = "aircon",
    "architecturalstyletypeid" = "architectural_style",
    "pooltypeid10" = "flag_spa",
    "pooltypeid2" = "flag_pool_spa",
    "pooltypeid7" = "flag_pool_tub",
    "fips"="county"))

# numerical variable
variable_numeric = c("area_firstfloor_finished",
    "area_base", "area_base_living",
    "area_garage",
    "area_live_finished",
    "area_liveperi_finished",
    "area_lot",
    "area_patio",
    "area_pool",
    "area_shed",
    "area_total_calc",
    "area_total_finished",
    "area_unknown",
    "tax_building",
    "tax_land",
    "tax_property",
    "tax_total",
    "latitude",

```

```

        "longitude")
# discrete
variable_discrete = c("num_75_bath",
    "num_bath",
    "num_bathroom",
    "num_bathroom_calc",
    "num_bedroom",
    "num_fireplace",
    "num_garage",
    "num_pool",
    "num_room",
    "num_story",
    "num_unit")
variable_binary = c("flag_fireplace",
    "flag_tub",
    "flag_spa",
    "flag_pool_spa",
    "flag_pool_tub",
    "tax_delinquency")

# categorical variable
variable_nominal = c("aircon",
    "architectural_style",
    "county",
    "deck",
    "framing",
    "heating",
    "id_parcel",
    "material",
    "region_city",
    "region_county",
    "region_neighbor",
    "region_zip",
    "story",
    "zoning_landuse",
    "zoning_landuse_county")
variable_ordinal = c("quality")

# date
variable_date = c("tax_year",
    "build_year",
    "tax_delinquency_year",
    "trans_year",
    "trans_month",
    "trans_day",
    "trans_date",
    "trans_weekday")

# others
variable_unstruct = c("zoning_property")

# don't understand
variable_unknown = c('censustractandblock',
    'rawcensustractandblock')

```

```

# Conversion
# - convert some binary to 0, 1
# - convert to date to int
# - convert to numeric to double
# - convert to discrete to int
# - convert to categorical to character

train[train$flag_fireplace == "", "flag_fireplace"] = 0
train[train$flag_fireplace == "true", "flag_fireplace"] = 1
train[train$flag_tub == "", "flag_tub"] = 0
train[train$flag_tub == "true", "flag_tub"] = 1
train[train$tax_delinquency == "", "tax_delinquency"] = 0
train[train$tax_delinquency == "Y", "tax_delinquency"] = 1
# convert to date to int
train[,variable_date] = sapply(train[,variable_date], as.character)
# convert to numeric to double
train[,variable_numeric] = sapply(train[,variable_numeric], as.numeric)
# convert to discrete to int
train[,c(variable_discrete, variable_binary)] = sapply(train[,c(variable_discrete, variable_binary)], as.integer)
# convert to categorical to character
train[,c(variable_nominal, variable_ordinal)] = sapply(train[,c(variable_nominal, variable_ordinal)], as.character)

#drop the redundant value
train <- train[, !(names(train) %in% c("num_bath","num_bathroom_calc","region_county","tax_year"))]

#deleting missing value
num.NA <- sort(colSums(sapply(train, is.na)))
remain.col <- names(num.NA)[which(num.NA <= 0.8 * dim(train)[1])] #
train <- train[,remain.col]
#####

#2. missing value imputation
check.na <- function(train){
  mis.col <- colSums(is.na(train))
  mis.col <- mis.col[mis.col>0]
  return(mis.col)
}

##(1) handle those have relative small missing value -> tax/total, tax_land/ tax_property

train[which(is.na(train$tax_total)),"tax_total"] <- train[which(is.na(train$tax_total)),"tax_property"]*
  quantile(train$tax_property/train$tax_total, 0.5, na.rm = T)

train[which(is.na(train$tax_property)),"tax_property"] <- train[which(is.na(train$tax_property)),"tax_total"]*
  quantile(train$tax_property/train$tax_total, 0.5, na.rm = T)

train[which(is.na(train$tax_land)),"tax_land"] <- train[which(is.na(train$tax_land)),"tax_total"]*
  quantile(train$tax_land/train$tax_total, 0.5, na.rm = T)

train[which(is.na(train$tax_building)),"tax_building"] <- train[which(is.na(train$tax_building)),"tax_total"]*
  quantile(train$tax_building/train$tax_total, 0.5, na.rm = T)
train[which(is.na(train$tax_building)),"tax_land"]
check.na(train)

```

```

#2. then impute region_zip/region_neighbor/region_city based on latitude/longitude
# actually we might just delete region_neighbor, because the missing value does not contain
#anyinformation, and we have to bother to impute it
with(train,t.test(logerror~ is.na(region_neighbor)))

# p-value is 0.07, we can accept the null hypothesis, region_neighbor does not contain info
#for those who are interested, take a look at how to find the nearest datapoint
#https://stackoverflow.com/questions/21977720/r-finding-closest-neighboring-point-and-number-of-neighbors
geo.df <- train[,c("latitude","longitude","region_zip","region_city")]
library(sp)
library(rgeos)
Impute.Zip<- function(row,search.range=1000){
  geo.df.trunc <-subset(geo.df, abs(geo.df$longitude + geo.df$latitude -
                                unlist(row["longitude"])-unlist(row["latitude"]))<search.range)

  coordinates(geo.df.trunc) <- ~longitude+latitude
  d <- gDistance(geo.df.trunc, byid=T)
  k = 2
  min.d <- apply(d, 1, function(x) order(x, decreasing=F)[k])
  neighbor <- min.d[rownames(row)]
  while(is.na(geo.df.trunc[neighbor,"region_zip"]$region_zip)){
    k = k + 1
    min.d <- apply(d, 1, function(x) order(x, decreasing=F)[k])
    neighbor <- min.d[rownames(row)]
  }
  return(geo.df.trunc[neighbor,"region_zip"]$region_zip)
}
na.zip.rowname <- which(is.na(train$region_zip))
for (i in na.zip.rowname){
  row <- geo.df[i,]
  geo.df[i,"region_zip"] = Impute.Zip(row)
}

#similarly, impute city
Impute.city<- function(row,search.range=1000){
  geo.df.trunc <-subset(geo.df, abs(geo.df$longitude + geo.df$latitude -
                                unlist(row["longitude"])-unlist(row["latitude"]))<search.range)

  coordinates(geo.df.trunc) <- ~longitude+latitude
  d <- gDistance(geo.df.trunc, byid=T)
  k = 2
  min.d <- apply(d, 1, function(x) order(x, decreasing=F)[k])
  neighbor <- min.d[rownames(row)]
  while(is.na(geo.df.trunc[neighbor,"region_city"]$region_city)){
    k = k + 1
    min.d <- apply(d, 1, function(x) order(x, decreasing=F)[k])
    neighbor <- min.d[rownames(row)]
  }
  return(geo.df.trunc[neighbor,"region_city"]$region_city)
}
na.city.rowname <- which(is.na(train$region_city))

```

```

count = 0
for (i in na.city.rowname){
  count = count + 1
  print(count)
  print(i)
  row <- geo.df[i,]
  geo.df[i,"region_city"] = Impute.city(row)
}

train[,c("region_zip","region_city")] = geo.df[,c("region_zip","region_city")]
train <- train[,!names(train) != "region_neighbor"]

#delete censustractandblock, region_neighbor -> do not have explain power
train <- train[,!names(train) %in%c("X.1","X","censustractandblock","region_neighbor")]
check.na(train)
names(train)
plot(train$area_lot,train$area_lot)
#redundant variable, delete area_live_finished
train <- train[,!names(train) == "area_live_finished"]

#fisrt handle some simple one, we create a new level to indicate missing value
train$aircon <- ifelse(is.na(train$aircon),"missing",train$aircon)
train$num_garage <- ifelse(is.na(train$num_garage),"missing",train$num_garage)
train$area_garage <- ifelse(is.na(train$area_garage),"missing",train$area_garage)
train$heating<- ifelse(is.na(train$heating),"missing",train$heating)
train$num_story<- ifelse(is.na(train$num_story),"missing",train$num_story)
train$quality<- ifelse(is.na(train$quality),"missing",train$quality)
train$build_year<- ifelse(is.na(train$build_year),"missing",train$build_year)
check.na(train)
# then let us impute the remaining missing value using library(mice)
# we have already got number information, for now, I did not see the explanation power in num_unit
train <- train[,!names(train) == "num_unit"]
library("mice")
#put all the columns related to area
area.col <- c("area_lot","area_total_calc","num_bedroom","num_bathroom","num_room","tax_total")
imputed_Data <- mice(train[,area.col], m=1, maxit = 10, method = 'pmm', seed = 500)
plot(train$area_lot,complete(imputed_Data,1)$area_lot)
train$area_lot = complete(imputed_Data,1)$area_lot
abline(0,1)

train$area_total_calc = complete(imputed_Data,1)$area_total_calc
#
check.na(train)
# here, numerical value can't contain NA
plot(table(train$num_garage))
#keep one, delete garage_area
train <- train[,!names(train) == "area_garage"]
#delete num_story
with(train, t.test(logerror~ is.na(num_story)))
train <- train[,!names(train) == "num_story"]
with(train, t.test(logerror~ is.na(num_garage)))
train <- train[,!names(train) == "num_garage"]

```

## 2 Feature Engineering

We separate all the features into several different categories, then perform the feature engineering within categories. Here is a summary of feature engineering.

### (1) Date-related features


- (a) “weekend”: whether the sold day is weekend
- (b) “Q2”: whether the sold day belongs to Q2
- (c) “ancient level”: categorize built year by cut (1940, 1995)

The above three features cover up day - week - season - year.

Note: the reasons I came up with the coming three features is due to visualization result by plot mean(error)~category.

```
variable_date = c("build_year",
                  "trans_year",
                  "trans_month",
                  "trans_day",
                  "trans_date",
                  "trans_weekday")

# trans_weekday
par(mfrow = c(1,1))
boxplot(subset(train,trans_weekday == "Monday" &
               abs(logerror)<0.145)$logerror,
        subset(train,trans_weekday == "Tuesday" &
               abs(logerror)<0.145)$logerror,
        subset(train,trans_weekday == "Wednesday" &
               abs(logerror)<0.145)$logerror,
        subset(train,trans_weekday == "Thursday" &
               abs(logerror)<0.145)$logerror,
        subset(train,trans_weekday == "Friday" &
               abs(logerror)<0.145)$logerror,
        subset(train,trans_weekday == "Saturday" &
               abs(logerror)<0.145)$logerror,
        subset(train,trans_weekday == "Sunday" &
               abs(logerror)<0.145)$logerror
        )
```



```
with(train,t.test(logerror~ trans_weekday == "Saturday"|trans_weekday == "Sunday"))

##
##  Welch Two Sample t-test
##
## data:  logerror by trans_weekday == "Saturday" | trans_weekday == "Sunday"
## t = 3.941, df = 1195.1, p-value = 8.586e-05
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
##  0.005499136 0.016402573
## sample estimates:
## mean in group FALSE mean in group TRUE
##      0.0115915048      0.0006406504

#new feature 1.
train$weekend =ifelse(train$trans_weekday== "Saturday"|train$trans_weekday == "Sunday",
                      1,0)

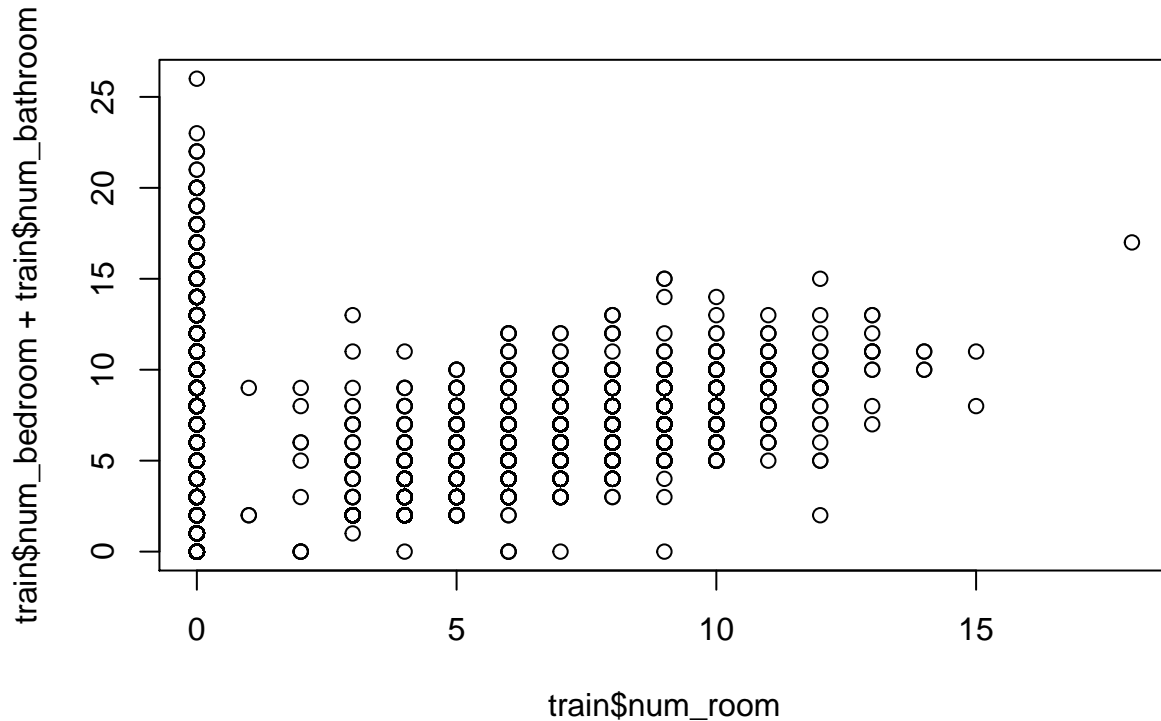
train$ancient.level <- ifelse(train$build_year == "missing","missing",
                              ifelse(as.integer(train$build_year)>1995,"new",
                                      ifelse(as.integer(train$build_year)>1940,"old","ancient"))))

train$Q2 = with(train,ifelse(trans_month %in% c("4","5","6"),"1","0"))
date.feature <- c("weekend","Q2","ancient.level")
```

(2) room/area related features



```
plot(train$num_room, train$num_bedroom + train$num_bathroom)
```



```
with(train,t.test(abs(logerror)~ num_room >= num_bedroom + num_bathroom))
```

```
##
## Welch Two Sample t-test
##
## data: abs(logerror) by num_room >= num_bedroom + num_bathroom
## t = 5.7024, df = 34396, p-value = 1.191e-08
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
## 0.004327230 0.008859982
## sample estimates:
## mean in group FALSE mean in group TRUE
## 0.06997593 0.06338233
```

```
train$right_room <- with(train,ifelse(num_room >= num_bedroom + num_bathroom,1,0))
train$bed_to_bath <- with(train,ifelse(num_bathroom == 0, 0,num_bedroom/num_bathroom))
cor(train$bed_to_bath,train$logerror)
```

```
## [1] -0.006295695
```

```
cor(train$bed_to_bath,abs(train$logerror))
```


```
## [1] 0.004323492
```

```
cor(train$num_bathroom,abs(train$logerror))
```

```
## [1] 0.001316901
train$are_per_room<- with(train,ifelse(num_room == 0, 0, area_total_calc/num_room))
train$zero_room <- with(train, ifelse(num_room ==0 ,1,0))
cor(train$are_per_room,abs(train$logerror))

## [1] -0.02167194
cor(train$area_total_calc,abs(train$logerror))

## [1] 0.03953064
par(mfrow = c(1,1))
with(train,plot(area_lot),plot(area_total_calc))
```



```
room.area.feature <- c("num_room","num_bedroom","num_bathroom","right_room","bed_to_bath",
  "are_per_room","zero_room","area_total_calc")
```

### (3) facility realted features

```
table(train$aircon)

##
##      1      11      13      3      5      9 missing
## 26668     63   1833      1    215      1   61494
by(train,train$aircon,function(x){mean(x$logerror)})
```

```

## train$aircon: 1
## [1] 0.01285241
## -----
## train$aircon: 11
## [1] 0.02629365
## -----
## train$aircon: 13
## [1] 0.01546552
## -----
## train$aircon: 3
## [1] 0.0917
## -----
## train$aircon: 5
## [1] 0.01568233
## -----
## train$aircon: 9
## [1] 0.01
## -----
## train$aircon: missing
## [1] 0.01070144

#rebuild level based on mean log_error
train$new.aircon = with(train,ifelse(aircon=="missing",aircon,
                                     ifelse(aircon=="5"|aircon=="13","5/13",
                                             ifelse(aircon=="11","11","1/3/9"))))

table(train$heating)

##
##      1      10      11      12      13      14      18      2      20
##      13      2      1      1      76      2      25 38303      97
##      24      6      7 missing
##    1071     970  15519  34195

by(train,train$heating,function(x){mean(x$logerror)})

## train$heating: 1
## [1] 0.02457692
## -----
## train$heating: 10
## [1] -0.0293
## -----
## train$heating: 11
## [1] -0.0151
## -----
## train$heating: 12
## [1] -0.0131
## -----
## train$heating: 13
## [1] -1.184211e-05
## -----
## train$heating: 14
## [1] -0.0049
## -----
## train$heating: 18

```

```
## [1] 0.027436
## -----
## train$heating: 2
## [1] 0.01324539
## -----
## train$heating: 20
## [1] -0.0006814433
## -----
## train$heating: 24
## [1] -0.01097722
## -----
## train$heating: 6
## [1] 0.009798144
## -----
## train$heating: 7
## [1] 0.007491887
## -----
## train$heating: missing
## [1] 0.01205165

train$new.heating = with(train,ifelse(heating == "missing",heating,
                                     ifelse(heating %in% c("10","11","12","13","14","20","24"),"neg",
                                               ifelse(heating %in% c("6","7"),"6/7","2/1/18"))))
train$missing.heating <- with(train,ifelse(heating == "missing",1,0))
train$missing.aircon <- with(train,ifelse(aircon == "missing",1,0))
facility.feature <- c("flag_tub","flag_fireplace","missing.aircon","new.heating")
```

#### (4) tax related features

```
with(train,cor(tax_delinquency,logerror))

## [1] 0.01893559

train$tax.perc <- train$tax_property/ train$tax_total
with(train,cor(tax.perc,logerror))

## [1] -0.003674848

train$right.tax <-ifelse(train$tax.perc < 0.1,1,0)
train$build.to.total <- train$tax_building/train$tax_total
with(train,cor(build.to.total,logerror))

## [1] 0.01948048

train$tax.per.area <- train$tax_property/train$area_total_calc
tax.feature <- c("tax_delinquency","tax_total","tax_property","tax_land","right.tax","tax.per.area","bu
```

#### (5)Overall evaluation feature(Collapse quality )

```
quality.error <- by(train,train$quality, function(x){mean(x$logerror)})
table(train$quality)

##
##      1      10      11      12      4      6      7      8 missing
```

```
##      2627      1461          1      119      23839          2      29310          5      32911
train$new.quality <- with(train,ifelse(quality=="missing",quality,
                                     ifelse(quality %in% c("12","7","11"),"12-7-11",
                                     ifelse(quality %in% c("4","1"), "4-1","6-8-10")
                                     )
))

quality.feature <- c("new.quality")
```


## (6) geometry related feature

```
zip.num <- by(train, train$region_zip, function(x) nrow(x))
train$zip.num <- zip.num[train$region_zip]
zip.area <- by(train, train$region_zip, function(x)
{abs((max(x$longitude) - min(x$longitude))*(max(x$latitude) - min(x$latitude))))})
train$zip.area <- zip.area[train$region_zip]

zip.density <- ifelse(zip.area == 0 , 0 ,zip.num/zip.area)
zip.density.accuracy <- ifelse(zip.num>50,1,0)
zip.density.revised <- zip.density*zip.density.accuracy
train$zip.density.revised <- zip.density.revised[train$region_zip]


plot(density(zip.density))
```

## density.default(x = zip.density)



N = 388 Bandwidth = 2.167e-08

```
plot(zip.density)
```



```
train$zip.density <- zip.density[train$region_zip]
with(train,cor(zip.density,logerror))
```

```
## [1] -0.01249644
```

```
with(train,cor(zip.density,abs(logerror)))
```

```
## [1] 0.0162063
```

```
train$zip.density <- zip.density[train$region_zip]
```

```
##
```

```
city.zip <- by(train,train$region_city, function(x){length(unique(x$region_zip))})
```

```
train$zip.per.city <- city.zip[train$region_city]
```

```
with(train,cor(zip.per.city ,logerror))
```

```
## [1] -0.006879891
```

```
city.area <- by(train, train$region_city, function(x){abs((max(x$longitude) - min(x$longitude))*(max(x$
```

```
city.num <- by(train, train$region_city, function(x) nrow(x))
```

```
city.density <- city.area/city.num
```

```
train$city.density <- city.density[train$region_city]
```

```
with(train,cor(city.density ,abs(logerror)))
```

```
## [1] 0.01449351
```

```
train$county6111 <- with(train, ifelse(county == "6111",1,0))
####for now.
geo.feature<- c("county6111","zip.density.revised")
```

## build linear model

Here, we do not list every linear model built. Instead, we just give an example to show the basic idea of how to choose the model using residual plot diagnosis.

```
selected.feature <- c("logerror",geo.feature,quality.feature,tax.feature,facility.feature,room.area.feature)
reg.df <- train[,selected.feature]
```

```
full.model <- lm(logerror~.,data = reg.df)
summary(full.model)
```

```
##
## Call:
## lm(formula = logerror ~ ., data = reg.df)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-4.6260	-0.0378	-0.0059	0.0275	4.7276


```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-5.332e-04	1.216e-02	-0.044	0.96503
county6111	1.418e-03	2.495e-03	0.568	0.56995
zip.density.revised	-1.010e+04	4.008e+03	-2.518	0.01179 *
new.quality4-1	-9.122e-04	1.780e-03	-0.512	0.60839
new.quality6-8-10	-2.195e-03	4.834e-03	-0.454	0.64985
new.qualitymissing	1.201e-03	2.944e-03	0.408	0.68324
tax_delinquency	2.450e-02	3.862e-03	6.344	2.25e-10 ***
tax_total	3.995e-08	5.688e-09	7.024	2.18e-12 ***
tax_property	-2.850e-06	2.738e-07	-10.411	< 2e-16 ***
tax_land	-1.866e-08	6.175e-09	-3.021	0.00252 **
right.tax	4.260e-02	6.262e-03	6.803	1.03e-11 ***
tax.per.area	-5.028e-05	7.097e-05	-0.708	0.47869
build.to.total	2.306e-04	3.680e-03	0.063	0.95003
tax.perc	-1.795e-04	1.232e-03	-0.146	0.88420
flag_tub	-1.518e-02	3.493e-03	-4.344	1.40e-05 ***
flag_fireplace	-1.030e-03	1.118e-02	-0.092	0.92662
missing.aircon	2.773e-03	1.597e-03	1.736	0.08253 .
new.heating6/7	-7.506e-04	2.098e-03	-0.358	0.72051
new.heatingmissing	-6.748e-03	2.801e-03	-2.409	0.01600 *
new.heatingneg	-2.652e-02	5.404e-03	-4.908	9.24e-07 ***
num_room	-1.133e-03	8.110e-04	-1.398	0.16225
num_bedroom	4.710e-04	1.184e-03	0.398	0.69080
num_bathroom	-2.737e-03	1.595e-03	-1.716	0.08615 .
right_room	-1.581e-02	5.451e-03	-2.900	0.00373 **
bed_to_bath	-2.943e-03	1.967e-03	-1.496	0.13457
are_per_room	-4.017e-05	1.372e-05	-2.929	0.00340 **
zero_room	-3.568e-02	8.755e-03	-4.076	4.59e-05 ***
area_total_calc	1.107e-05	1.245e-06	8.895	< 2e-16 ***

```
## weekend          -1.229e-02  4.866e-03 -2.526  0.01153 *
## Q21             -6.899e-03  1.135e-03 -6.077  1.23e-09 ***
## ancient.levelmissing 1.148e-02  7.586e-03  1.514  0.13009
## ancient.levelnew    4.555e-03  2.508e-03  1.816  0.06935 .
## ancient.levelold    3.468e-03  1.891e-03  1.834  0.06661 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1606 on 90242 degrees of freedom
## Multiple R-squared:  0.006293,    Adjusted R-squared:  0.005941
## F-statistic: 17.86 on 32 and 90242 DF,  p-value: < 2.2e-16
```

Then let us start the model selection using lasso regression.

```
library(glmnet)
ind = model.matrix(~.,reg.df[,-1])
dep = reg.df$logerror
fit <- glmnet(x=ind, y=dep)
plot(fit, xvar = "lambda", label = T)
```

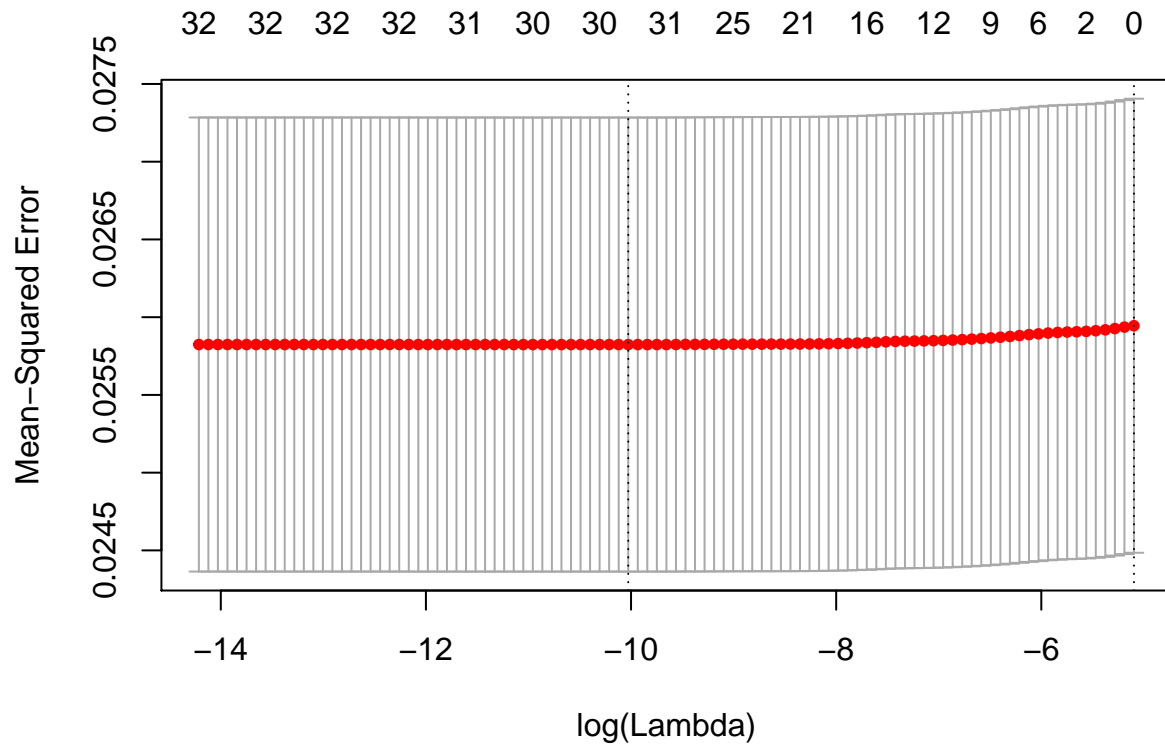


```
cvfit <- cv.glmnet(ind, dep)
cvfit$lambda.min
```

```
## [1] 4.417089e-05
```



```
plot(cvfit)
```



```
x = coef(cvfit, s = "lambda.min")
x
```

```
## 34 x 1 sparse Matrix of class "dgCMatrix"
##                               1
## (Intercept)                -2.301513e-02
## (Intercept)                  .
## county6111                   1.611377e-03
## zip.density.revised          -9.806271e+03
## new.quality4-1               -2.177119e-04
## new.quality6-8-10            -4.056155e-04
## new.qualitymissing           1.530044e-03
## tax_delinquency              2.402104e-02
## tax_total                    3.359265e-08
## tax_property                 -2.626326e-06
## tax_land                     -1.367046e-08
## right.tax                     4.424596e-02
## tax.per.area                 -5.256054e-05
## build.to.total               1.140127e-03
## tax.perc                     -8.133288e-05
## flag_tub                     -1.431587e-02
## flag_fireplace               .
## missing.aircon               2.351777e-03
## new.heating6/7              -3.513169e-04
```

```
## new.heatingmissing -5.612047e-03
## new.heatingneg -2.357419e-02
## num_room -2.902971e-04
## num_bedroom -6.866857e-05
## num_bathroom -1.204989e-03
## right_room -8.081183e-03
## bed_to_bath -1.366878e-03
## are_per_room -2.402737e-05
## zero_room -1.769374e-02
## area_total_calc 1.013892e-05
## weekend -1.191319e-02
## Q21 -6.793608e-03
## ancient.levelmissing 6.902538e-03
## ancient.levelnew 3.975123e-03
## ancient.levelold 3.008395e-03
```

We can see lasso regression did not really help us choose some features.

However, we want to use residual plot to seek for new opportunities to improve the model. I will just give an example for the feature bedroom to bathroom rate.

After performing the square root transformation, the residual plot will look better. And we trained the model again, the adjusted R-square does improve as well.

Then, for each feature, we tried to perform the some transformation, also include some intersection terms.

After trying different transformations and interactions, I showed the best result below.

```
mod <- lm(logerror~.-zip.density.revised +I((train$zip.density.revised)^(1/18))
  -are_per_room + I(are_per_room^(1/15))
  -county6111-new.quality
  #-tax_land +I(tax_land^(1/2))
  -tax_property + I(tax_property^(1/60))
  -num_room
  -bed_to_bath + I((bed_to_bath)^(1/2))
  -tax.per.area + I(tax.per.area^(1/50))
  -flag_fireplace
  -missing.aircon
  -tax.perc + I(tax.perc^(1/12))
  -right.tax
  + right.tax:I(tax_land^(1/3))
  +I(build.to.total^(1/60))
  + flag_tub : new.heating
  +right_room : zero_room
  -flag_tub
  +Q2:weekend

  ,data = reg.df)
summary(mod)
```

```
##
## Call:
## lm(formula = logerror ~ . - zip.density.revised + I((train$zip.density.revised)^(1/18)) -
## are_per_room + I(are_per_room^(1/15)) - county6111 - new.quality -
## tax_property + I(tax_property^(1/60)) - num_room - bed_to_bath +
## I((bed_to_bath)^(1/2)) - tax.per.area + I(tax.per.area^(1/50)) -
## flag_fireplace - missing.aircon - tax.perc + I(tax.perc^(1/12)) -
```

```

##      right.tax + right.tax:I(tax_land^(1/3)) + I(build.to.total^(1/60)) +
##      flag_tub:new.heating + right_room:zero_room - flag_tub +
##      Q2:weekend, data = reg.df)
##
## Residuals:
##      Min        1Q    Median        3Q        Max
## -4.6221 -0.0381 -0.0058  0.0284  4.7295
##
## Coefficients:
##                                Estimate Std. Error t value
## (Intercept)                   1.847e+00  1.420e-01  13.009
## tax_delinquency                2.456e-02  3.856e-03   6.371
## tax_total                      2.173e-08  5.454e-09   3.984
## tax_land                      -3.180e-08  6.898e-09  -4.610
## build.to.total                 -2.168e-02  7.113e-03  -3.049
## new.heating6/7                 -1.276e-03  1.814e-03  -0.703
## new.heatingmissing             -5.660e-03  1.620e-03  -3.494
## new.heatingneg                 -2.974e-02  4.939e-03  -6.022
## num_bedroom                    4.983e-03  1.267e-03   3.933
## num_bathroom                  -6.229e-03  1.686e-03  -3.695
## right_room                     -1.496e-02  6.528e-03  -2.291
## zero_room                     -1.878e-01  8.083e-02  -2.324
## area_total_calc                1.732e-05  2.058e-06   8.416
## weekend                        -1.243e-02  5.171e-03  -2.404
## Q21                           -6.639e-03  1.136e-03  -5.842
## ancient.levelmissing           3.348e-02  9.264e-03   3.614
## ancient.levelnew               1.269e-02  2.381e-03   5.329
## ancient.levelold              5.460e-03  1.813e-03   3.012
## I((train$zip.density.revised)^(1/18)) -4.459e-02  1.171e-02  -3.808
## I(are_per_room^(1/15))         -1.185e-01  5.378e-02  -2.203
## I(tax_property^(1/60))         -1.798e+00  2.325e-01  -7.734
## I((bed_to_bath)^(1/2))         -2.377e-02  5.328e-03  -4.462
## I(tax.per.area^(1/50))         4.524e-01  1.940e-01   2.333
## I(tax.perc^(1/12))             -1.625e-01  5.241e-02  -3.100
## I(build.to.total^(1/60))       7.430e-02  1.073e-02   6.923
## right.tax:I(tax_land^(1/3))    3.703e-04  1.587e-04   2.333
## flag_tub:new.heating2/1/18     1.678e-03  6.872e-03   0.244
## flag_tub:new.heating6/7       -8.010e-03  1.422e-02  -0.563
## flag_tub:new.heatingmissing    -1.952e-02  4.145e-03  -4.708
## flag_tub:new.heatingneg        4.665e-02  1.135e-01   0.411
## right_room:zero_room          -1.593e-02  1.151e-02  -1.384
## weekend:Q21                     5.243e-03  1.506e-02   0.348
##                                Pr(>|t|)
## (Intercept)                   < 2e-16 ***
## tax_delinquency                1.89e-10 ***
## tax_total                      6.78e-05 ***
## tax_land                      4.04e-06 ***
## build.to.total                 0.002299 **
## new.heating6/7                 0.481892
## new.heatingmissing             0.000475 ***
## new.heatingneg                 1.73e-09 ***
## num_bedroom                    8.39e-05 ***
## num_bathroom                   0.000220 ***
## right_room                     0.021942 *

```

```

## zero_room          0.020149 *
## area_total_calc    < 2e-16 ***
## weekend              0.016205 *
## Q21                 5.16e-09 ***
## ancient.levelmissing 0.000301 ***
## ancient.levelnew    9.89e-08 ***
## ancient.levelold    0.002594 **
## I((train$zip.density.revised)^(1/18)) 0.000140 ***
## I(are_per_room^(1/15)) 0.027566 *
## I(tax_property^(1/60)) 1.05e-14 ***
## I((bed_to_bath)^(1/2)) 8.15e-06 ***
## I(tax.per.area^(1/50)) 0.019670 *
## I(tax.perc^(1/12)) 0.001937 **
## I(build.to.total^(1/60)) 4.46e-12 ***
## right.tax:I(tax_land^(1/3)) 0.019644 *
## flag_tub:new.heating2/1/18 0.807115
## flag_tub:new.heating6/7 0.573320
## flag_tub:new.heatingmissing 2.50e-06 ***
## flag_tub:new.heatingneg 0.680953
## right_room:zero_room 0.166344
## weekend:Q21          0.727687
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1603 on 90243 degrees of freedom
## Multiple R-squared:  0.01029,    Adjusted R-squared:  0.009954
## F-statistic: 30.28 on 31 and 90243 DF,  p-value: < 2.2e-16

```